

Rahul Raj Singh

(240)-706-4275 | rahulsingh061491@gmail.com | www.linkedin.com/in/rahmesses1423

PROFESSIONAL SUMMARY

Senior Data Engineer with 7+ years of experience building analytics-ready data platforms on Snowflake, dbt, and AWS. Expertise in Medallion architectures, semantic and dimensional modeling, and data quality frameworks that transform complex, high-volume data into trusted assets. Proven track record of translating disparate source data into actionable business insights by partnering with Finance and Product teams to deliver reliable, high-impact data products.

CORE SKILLS

Analytics Engineering & Modeling: dbt, Medallion Architecture, Semantic & Dimensional Modeling, SCD1/SCD2

Data Platforms & Languages: Snowflake, AWS (S3, RedShift, Glue, Lambda, DMS, Kinesis), Apache Kafka, Python, SQL, PySpark

Data Quality & Governance: PyTest, Anomalo, Datadog, Data Versioning, Data Contracts, Schema Evolution, GDPR, PII Masking

Orchestration & Infrastructure: Airflow, Apache NiFi, Terraform, Git, Docker, CI/CD, Technical Design Reviews

BI, Analytics & AI: Tableau, Metric Governance, CursorAI, LLM-Assisted Logical Refactoring, Automated Documentation

PROFESSIONAL EXPERIENCE

Senior Data Engineer | The Walt Disney Company | Baltimore, MD

May 2022 – Present

- Led the design and implementation of a modern **Medallion-based ELT platform**, processing **100M+** records daily, improving data reliability while enhancing BI user trust by **~90 %** and reducing developer maintenance overhead by **~65%**.
- Democratized data access by designing a governed **semantic layer (150+ data models)** in **partnership** with Finance and Product teams; enabled seamless **self-service BI** and reduced ad-hoc reporting by **~75%**.
- Established **dbt modeling** standards by engineering reusable **macros** and templates for incremental processing and versioning; ensured long-term maintainability for a catalog of **500+ data assets**.
- Architected an **event-driven** pipeline for **15–16M daily Kafka streaming events** using Snowflake native orchestration (Pipes, Streams, Tasks, Stored Procedures); achieved **99.9% uptime** and reduced manual troubleshooting by **70%**.
- Designed a '**Quality-as-Code**' observability framework with **PyTest** for developers, later extending to **Anomalo** and **Datadog** to monitor **Critical Data Elements**; reduced manual triage by **~85%** while ensuring **24/7 reliability** for production assets.
- Lead ongoing Snowflake performance and **cost optimization** efforts through **clustering** and model-level refactoring, with **~80%** compute savings (**~\$45K/year**) and **~82%** runtime reductions (**~60m → ~10m**) achieved to date across **500M+ row** datasets.
- Provide **technical leadership** to a team of **4+ engineers**, defining modeling standards, conducting rigorous design reviews, and aligning cross-functional stakeholders to deliver reliable, **production-grade** data products.

Associate Manager – Data Engineering | Factspace Inc. | Baltimore, MD

May 2021 – May 2022

- Led 3 parallel data engineering teams (**scaling from 2 to 12 engineers**) as the onsite engagement lead; owned end-to-end prioritization and delivery for ETL, legacy system sustainment, and global cloud migration projects.
- Co-architected and delivered an AWS-based enterprise data lake, migrating **250+ tables** from **10+** on-prem and cloud data sources to establish a unified foundation for downstream analytics and reporting.

Senior Associate – Data Engineering | Factspace Inc. | Baltimore, MD

Apr 2020 – Apr 2021

- Designed and implemented hybrid batch and streaming pipelines using **AWS Glue, Apache NiFi, and PySpark** to process **over 5M behavioral events per day**, supporting production recommender systems.
- Co-designed an **AWS Redshift** ingestion framework organizing raw, historical, syndicated, and curated layers to deliver feature-ready datasets for high-performance data science models.

Software Engineering Analyst | Accenture | Bangalore, India

Oct 2015 – June 2018

- Engineered a **Python-based** framework to automate ingestion of **10M+** street-view images of broadband network assets from Google Maps API, replacing manual browser-based inspections resulting in **\$2.3M in client cost savings**.
- Developed a **ML-model** using **TensorFlow** and **OpenCV** to automate the detection of network assets. Achieved a **6x** improvement in execution speed and a **90% reduction** in manual processing time.

EDUCATION

University of Maryland, Robert H. Smith School of Business - M.S. in Information Systems

Aug 2018 – Dec 2019

Visvesvaraya Technological University, Siddaganga Institute of Technology - B.E. in Mechanical Engineering

Aug 2010 – Jun 2015